



# Deep Learning Lab

Introduction to the Lab activities

Master Big Data 2021

Andrea Cossu

→ [andrea.cossu@sns.it](mailto:andrea.cossu@sns.it)

→ <https://andreacossu.github.io/>

Ph.D. student in Data Science

@ Scuola Normale Superiore

@ University of Pisa

working on *Continual Learning*

Member of [PAI](#) lab and [CIML](#) group

[ContinualAI](#) Team Lead



# Calendar



- Monday, April 19th, 14-18
- Monday, April 26th, 14-18
- Tuesday, May 4th, 14-18
- Monday, May 10th, 14-18
- Monday, May 17th, 14-18
- Microsoft Teams, DL channel (+ separate rooms when needed)

# DL Lab tools



- Keras
  - + Tensorflow
- Colab
- Matplotlib, numpy, scikit-learn...



# DL Lab outline



- Practical Deep Learning with Keras (and Tensorflow)
- Deep Learning areas
  - Tabular data
  - Computer Vision
  - Sequential data processing (audio, natural language...)
  - Generative models
  - Graph structured data
- You will see more stuff than you actually need for the project!

# DL Lab objectives



- Understand the main *principles* behind a Deep Learning library
  - learn one, master all
- Autonomously navigate a deep learning library
  - first google it, then ask it
- Use a deep learning library to solve a specific problem/application
  - aka, fun (and/or \$\$\$)
- We will go **slowly**!  
Most of the things you need are already implemented :)

# Final Project



- Verify that you acquired the objectives (theory + lab)
- Address a new problem from a known area **your own way**
  - you can use whatever code you like
  - you have to understand it, first
  - so that you can justify your design choice (why this? Have you tried that?)



# Practical DL


(before digging deeper into the code)




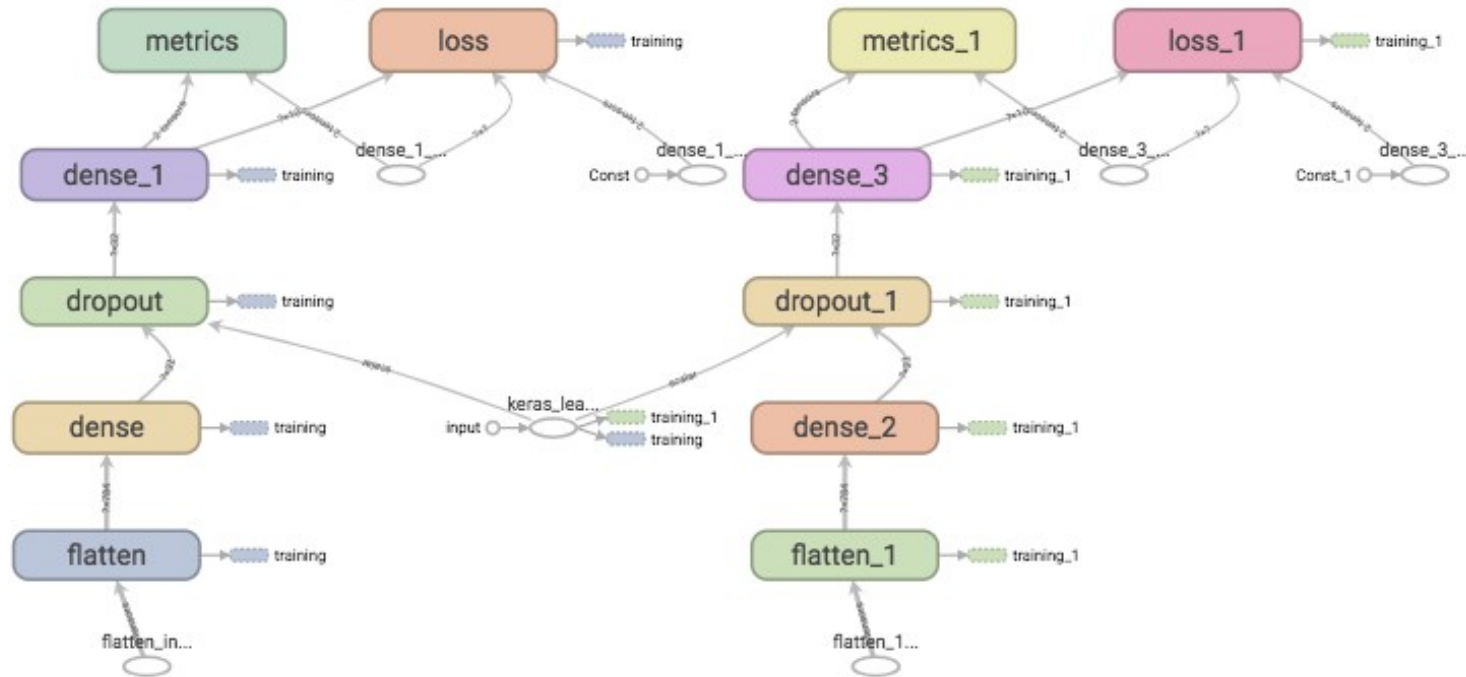
# Tensors (the computer science way)



- Multi-dimensional array filled with homogenous values
- They have (at least)
  - A shape/size of n dimensions  
( $D_1 \times D_2 \times \dots \times D_n$ )
  - A (data)type: float32, int32, bool...
- Manipulate tensors with common math operators



# Computational Graph



# Tensorflow



- Low level API to build DL models, ML pipelines...
- Deal directly with tensors, build models by nesting / composing different modules (e.g. layers)
- Manipulate gradients, control single variables
- Powerful, but more complex

# Keras



- High-level API to build DL models, ML pipelines...
- Deal directly with entire modules (e.g. layers)
- Fast prototyping of common paradigms (training loops, dataset manipulation...)
- Powerful and easier...
  - Unless you have to build specific DL models (e.g. change training loop requires mixing Keras + Tensorflow)
- Sometimes, high-level means harder to debug. Each function performs lots of computation, which is hidden to the user.

# PyTorch



- Tensorflow equivalent
  - better organized documentation
  - Less oriented to production – but it is changing fast
- High level API are available (e.g. FastAI, Lightning...)
- You don't want to mix the two in the beginning

# Monitoring tools



- You should always have a look at *how* your training is progressing
- Needed to understand which DL tricks to apply to improve performances
- Tensorboard (shipped with Keras), Weights and Biases (free for open-source projects), many others.



Let's code, shall we?